

NVIDIA (NVDA)

The AI Factory — a Dual Long & Short Case

SENIOR CALL: LONG-biased, two-sided — accumulate $\leq$ \$195, do not chase at \$210

Spot \$209.91 · reported P/E ~32x · ~\$5T market cap · 3-year horizon

The whole case turns on ONE number: NVDA's retained share of the AI-accelerator TAM

From Chip to Coherent Rack

The enabling shift — why this matters

- ▶ A frontier AI model is too big for any single chip — the bottleneck is interconnect, not raw FLOPS.
- ▶ NVLink + NVSwitch fuse 72 GPUs into one ~130 TB/s fabric, so a trillion-parameter model behaves as one accelerator. [ATTRIBUTED]
- ▶ The unit of competition moved from the chip to the rack — the “AI factory.”
- ▶ The sellable SKU moved up: ~\$30K GPU → ~\$3M GB200 NVL72 rack → ~\$8.8M Rubin rack. [ATTRIBUTED]

ELI5: 72 GPUs Wired to Act as One

The moat in plain English

- ▶ Imagine 72 office workers who, instead of emailing, share ONE desk and ONE whiteboard at conversational speed.
- ▶ That shared whiteboard (NVLink) is the moat.
- ▶ A rival can build a faster single worker (a custom ASIC) — but not the 72-at-one-desk room.
- ▶ ...and not the 18 years of software (CUDA) everyone already wrote for NVIDIA's room.

The Moat = Three Stacked Layers

The engineering choice that wins — a rival must beat all three at once

1 · SOFTWARE

~18 years of CUDA

cuDNN / TensorRT / NCCL

switching cost =

the whole ML stack

2 · INTERCONNECT

NVLink / NVSwitch

rack-scale coherence

72 GPUs as one

~130 TB/s domain

3 · PACKAGING

pre-booked TSMC

CoWoS-L allocation

(~60% of 2026)


the physical throttle

Annual cadence (Hopper → Blackwell → Rubin → Feynman) on R&D of only ~8.6% of sales

The Market: \$140B → over \$440B

AI-accelerator TAM — the end market compounds

- ▶ 2024 ~\$140B → 2030 >\$440B, ~25% CAGR. [ATTRIBUTED]
- ▶ Driven by hyperscaler capex (~\$725B in 2026, +~77%). [ATTRIBUTED]
- ▶ Sovereign AI + enterprise inference add pull.
- ▶ Even with share erosion, the dollar pie expands.



NVIDIA by the Numbers (FY2026)

Best-in-class — software economics on a hardware bill-of-materials

\$215.9B

Revenue (+65.5% YoY) [FACT]

71.1% / 60.4%

Gross / operating margin [FACT]

\$120.1B


Net income; \$102.7B OCF [FACT]

~76% ROE

Net cash; 2nd-cheapest P/E in cohort

Value Chain: Upstream → NVDA → Downstream

NVDA owns the highest-margin link




Risk: the same downstream customers are building their own silicon (TPU/Trainium/Maia/MTIA)

PxQ: The SKU Moved Up the Value Chain

Price × Quantity — supply-gated units, exploding ASP


- ▶ P: \$30K GPU → \$3M NVL72 rack → ~\$8.8M Rubin rack.
- ▶ Q: supply-gated by CoWoS-L + HBM capacity.
- ▶ Revenue: \$27B → \$61B → \$130B → \$216B.
[FACT]
- ▶ ~\$85B of incremental revenue in FY26 alone.



Competition Ranked by Threat

Where value could migrate — and what holds each rival back

- ▶ Custom ASIC (TPU/Trainium/Maia/MTIA): inference TCO 40-65% lower. [ATTRIBUTED]
- ▶ Broadcom & Marvell: merchant enablers of in-house silicon.
- ▶ AMD MI3xx: closing on price, behind on software/scale-up.
- ▶ NVDA still has the highest growth & margin in the group.



Commercial Scaling: Spec → Profit

Why the engineering choice drives margin, share and value

THE ENGINEERING CHOICE

THE BUSINESS OUTCOME

Sell the rack/system, not the die

ASP explodes (\$30K → \$3M rack)

CUDA lock-in

71% GM / 60% op margin [FACT]

Supply-gated (CoWoS-L/HBM)

No discounting; margin holds

Full-stack moat

ROIC far above WACC — value creation

The Bear Case — 4 Reasons to Be Careful

An honest short: a margin-and-share story, not a fraud short

1 · MARGIN ROLLED OVER

Gross margin 75.0% → 71.1%

at the TOP of the cycle [FACT]

HBM4/CoWoS-L cost + mix

2 · INFERENCE SHARE LEAK

inference ~2/3 of compute

ASICs 40-65% lower TCO

Anthropic→~1M TPUs [ATTRIBUTED]

3 · FY23 FAILURE MODE

revenue went flat (+0.2%)

op margin 37%→16% in 1yr [FACT]

now on a ~\$5.1T base

4 · CIRCULAR FINANCING

>\$40B AI-equity commits

vs \$10.6B cash [FACT]

Nortel/Lucent echoes

Are You Buying Near the Top?


The cycle & margin history — re-underwrite each cadence

- ▶ Gross-margin path: 64.9 → 56.9 → 72.7 → 75.0 → 71.1% [FACT] — peaked, now normalizing.
- ▶ Revenue compounded ~8x in three years with margins expanding, then rolling.
- ▶ The multiple embeds a lot on a ~\$5T base (32x fwd / ~43x trailing GAAP).
- ▶ A capex-cycle pause is a BUYING opportunity; a structural share break is the thesis-breaker.

Sensitivity Matrix: the Single Driver

Retained share × 2029 TAM → 3-yr price / IRR from \$209.91


Base cell
75% / \$440B
= \$244
+5% IRR



Bottom-left =
the short's
quadrant
(below today)

Valuation & Thesis: Bear / Base / Bull

Two trades on the same facts — the edge is the entry



- ▶ LONG: accumulate $\leq$ \$195, full $\leq$ \$180; ~3-yr hold.
- ▶ SHORT: put spread + AVGO/ASIC pair on a share-loss tell; target \$140-160.
- ▶ Base \$244-283 (+5-10%/yr).
- ▶ Bull \$346-372; Bear \$124-142.

Investment Philosophy

The open question, answered

“Find the one technical variable the whole valuation hangs on, price the tails, and let the ENTRY — not the narrative — carry the edge. Underwrite long AND short on the same facts; stress the load-bearing claims until one side refutes rather than merely rebuts; then size to the asymmetry. Express the bear as defined-risk relative value, not a conviction zero.”

What Would Change the Thesis

Intellectual honesty — the tells to watch

- ▶ Does the frontier stay NVDA-anchored, or does large-MoE training move to TPU/Trainium?
- ▶ Where does gross margin settle after the Rubin ramp — above or below ~67% structurally?
- ▶ Does hyperscaler capex ROI prove out, or pause off the ~\$725B 2026 budget?
- ▶ Power, not chips, as the ultimate ceiling on deliverable units (RAND).